

Research Paper Radar

Fresh arXiv research translated into systems, experiments, publication angles, and people worth knowing.

10	6	3	7
selected papers	verified codebases	deep reads today	experiment/publication ideas

EDITORIAL THESIS

Today's strongest cluster is not 'bigger models.' It is explicit control over how agents execute, test, remember, revise, and judge. Procedural Graphs externalizes what-to-do knowledge; ExecCritic separates verification from repair; MeClear and Revoked but Still Authoritative expose memory governance as an enforceable systems problem. PlannerForge shows how those ideas become a simulation-and-testing pipeline.

Evidence note

Eight selections are fresh arXiv preprints and should be treated as preliminary. PlannerForge reports EMNLP 2026 acceptance; the graph-memory survey reports ICKG 2026 acceptance. Numerical results are paper claims, not independent replications. Affiliations were checked against paper PDFs; code links were accepted only when the paper named the repository.

Executive brief // the three deep reads

#	Paper	Why today	First action
1	Procedural Graphs	Highest architecture fit: graph-guided execution plus validated self-evolution.	Read method + ablations; sketch a local task graph.
2	ExecCritic	Strongest reproducible coding-agent result and fail-closed systems lesson.	Clone repo; run one small repository task.
3	PlannerForge	Best bridge from agent pipelines to simulation, testing, robotics, and GIS.	Inspect scenario schema; map to a logistics domain.

Emerging research signals

Signal	Evidence in today's set	AutoNateAI implication
Control is becoming explicit	Graphs, frozen tests, agent-level interventions	Represent procedures, evaluators, and change histories as first-class artifacts.
Memory needs revocation semantics	MeClear + revocation-enforcement study + graph-memory survey	Build authorization and temporal validity into retrieval—not only prompts.
Evaluation is the product boundary	ExecCritic, SAEScientist, safety-judge wrappers	Independent evaluators need their own adversarial tests and failure budgets.
Agents are entering simulation toolchains	PlannerForge	Use agents to generate, mutate, execute, and score scenarios in other technical domains.
Theory may unlock new data products	Sheaves/narratives	Explore temporal invariants for evolving GIS, system, and relationship graphs.

Recommended allocation

Spend 60 minutes on Procedural Graphs, 45 minutes on ExecCritic plus its repository, and 45 minutes on PlannerForge plus scenario files. Use the remaining research block to pair MeClear with the revocation attack rather than reading them in isolation.

Search and category coverage

Primary scan: arXiv recent computer-science submissions for 8–9 September 2026, followed by paper-PDF verification, code-link extraction, and targeted checks of authoritative author/lab pages. Attention signals were used only for discovery, never as evidence of findings.

Domain	Selected coverage	Notes
Agents / multi-agent	Procedural Graphs; AgentGrad; MeClear	Execution, optimization, memory
Software engineering / devtools	ExecCritic	Independent test/repair scaffold
Robotics / simulation	PlannerForge	Scenario generation through benchmarking
AI evaluation / safety	SAEScientist; Style Over Substance	Research-agent and judge robustness
Data / graphs / memory	Graph Memory; Narratives; Revocation	Temporal, personalized, governed memory
Cybersecurity	MeClear; Revocation; safety judges	Integrity and policy enforcement
GIS / distributed systems / PL / HCI	No paper cleared today's bar	Continue daily search; no padding

Topic concentration



Counts are multi-label and show the edition's concentration, not field-wide publication volume.

Top papers leaderboard

Relevance score is an editorial AutoNateAI fit score, not paper quality or scientific validity.

#	Paper	Categories	Code	Fit
1	Procedural Graphs: Self-Evolving Execution Structures for LLM Agents	cs.AI / cs.CL / cs.MA	No verified repository	98
2	ExecCritic: Learn to Test, Test to Improve for Coding Agents	cs.AI / cs.CL / cs.SE	Repo	97
3	PlannerForge: LLM Agents for Scenario-Based Testing of Motion Planners	cs.AI / cs.CL / cs.RO	Repo	95
4	MeClear: Cooperative Game-Theoretic Attribution and Risk-Aware Memory Clearance	cs.AI / cs.SE	Repo	93
5	AgentGrad: Intervention-guided Prompt Optimization for Multi Agent Systems	cs.AI	No verified repository	91
6	SAEScientist-Bench: Can AI Agents Conduct Autonomous SAE Interpretability Research?	cs.AI / cs.CL / cs.LG	Repo	89
7	Graph-Based Personalized Memory for LLM Agents	cs.AI	No paper-specific repository	88
8	Time-Varying Data as Sheaves: an Invitation to Narratives	cs.AI / cs.MA / eess.SY / math.CT	No verified repository	86
9	Style Over Substance: Content-Invariant Wrappers Flip LLM Safety-Judge Verdicts	cs.AI	Repo	85
10	Revoked but Still Authoritative: Revocation Enforcement in Agent-Memory Systems	cs.AI / cs.CR	Repo	84

Selection logic

Ranking weights practical alignment, novelty of the system idea, code availability, evidence clarity, and experimentability. Fame and model scale receive no automatic preference. A paper can rank highly without code when it articulates a uniquely useful architecture, as Procedural Graphs does today.


Architecture intelligence

Procedural Graphs // editable execution control




Interpretation: the core move is not rigid workflow execution; it is a structured procedural prior that still lets the solver reason. The refiner only commits edits that preserve or improve held-out performance.

ExecCritic // independent behavioral verification



Interpretation: independent role training matters because the same trajectory can otherwise write a patch and a mutually agreeable but wrong test.

Research-to-publication operating loop



Paper briefs // 01-02

#1 Procedural Graphs: Self-Evolving Execution Structures for LLM Agents

98

Submitted 8 Sep 2026 | [cs.AI](#) / [cs.CL](#) / [cs.MA](#) | [PDF](#)

FIT

Authors	Yuxing Lu, Yicheng Chen, Shanchan Wu, Sercan Ö. Arık
Affiliations	Google; Georgia Tech; Peking University
Central question	Can long-horizon agents externalize procedural knowledge in an editable graph that improves from execution outcomes?
Method	Localize the active graph node; translate its neighborhood into situational guidance; contrast failed and successful trajectories; validate topology and attribute edits before committing them.
Key findings — paper claim	Paper reports consistent gains over memory baselines across multiple datasets, tasks, and LLMs; self-evolution further improves results and can repair flawed expert priors.
Limitations	No public code was found today. Guidance generation adds a model call; authors flag transfer and selective/reused guidance as future work.
Why it matters — interpretation	Closest match to Nathan's process graphs, agent orchestration, evaluation loops, and evolving system architecture.
Code / project	No verified repository

#2 ExecCritic: Learn to Test, Test to Improve for Coding Agents

97

Submitted 8 Sep 2026 | [cs.AI](#) / [cs.CL](#) / [cs.SE](#) | [PDF](#)

FIT

Authors	Leitian Tao et al.
Affiliations	UW-Madison; Microsoft Research; Georgia Tech
Central question	How can execution feedback improve coding agents without letting patch and test errors agree with each other?
Method	Use independent Test and Repair agents; qualify and freeze repository-native tests in a fail-closed harness; post-train each Qwen-3.5-35B-A3B role separately.
Key findings — paper claim	On SWE-bench Verified, the composed agents reach 72.6%, +11.4 points over the 61.2% no-test baseline. Weak generated tests can hurt: the base Test agent reduces resolution to 57.3%.
Limitations	Evaluation depends on isolated sandboxes and substantial compute. The paper notes test correctness, stronger repair learning, and broader evaluation as open directions.
Why it matters — interpretation	Immediately operationalizable for AutoNateAI coding agents, CI gates, repair evaluation, and customer-facing build sprints.
Code / project	Open repository

Paper briefs // 03-04

#3 PlannerForge: LLM Agents for Scenario-Based Testing of Motion Planners

95

Submitted 8 Sep 2026 | cs.AI / cs.CL / cs.RO | PDF

FIT

Authors	Yuan Gao et al.
Affiliations	TUM AVS / MIRMI; University College London
Central question	Can agents unify scenario generation, retrieval, modification, assessment, and benchmarking for autonomous-driving planners?
Method	Agent pipeline evaluated with 10 LLMs and five prompting conditions; tasks cover executable scenario generation, semantic retrieval, valid edits, and planner benchmarking.
Key findings — paper claim	Reports 193/200 executable scenarios versus 144 for a baseline, 92% rank-1 retrieval versus 67.5% BM25, and at N=400 a planner success increase from 50.4% to 70.2% with collisions falling from 19.0% to 8.4%.
Limitations	Automotive stack and scenario ontology constrain transfer; authors call for broader planners, simulators, and real-world validation. Accepted at EMNLP 2026.
Why it matters — interpretation	Strong bridge to Nathan's simulation, robotics, testing, and GIS/logistics interests; official code enables adaptation.
Code / project	Open repository

#4 MeClear: Cooperative Game-Theoretic Attribution and Risk-Aware Memory Clearance

93

Submitted 8 Sep 2026 | cs.AI / cs.SE | PDF

FIT

Authors	Boyu Yang, Jiazheng Sun, Zilong Lu, Zhi Qiu, Xin Peng, Jun Zheng
Affiliations	Fudan University; Beijing Institute of Technology
Central question	Which memories should a long-horizon agent suppress for a task when harmful effects are redundant or cooperative?
Method	Leave-one-out screening plus sampled Shapley attribution estimates joint harm; clearance is query scoped and avoids deleting persistent memory.
Key findings — paper claim	Across ten long-dialogue pools, reports 85.9% target recall and 82.3% task recovery, 25.5 points above a leave-one-out baseline.
Limitations	Game-theoretic estimation adds cost; evaluations use curated dialogue pools and do not prove durable privacy or revocation guarantees.
Why it matters — interpretation	Concrete design for safer persistent memory, especially when combined with adversarial revocation tests.
Code / project	Open repository

Paper briefs // 05-06

#5 AgentGrad: Intervention-guided Prompt Optimization for Multi Agent Systems

91

Submitted 8 Sep 2026 | cs.AI | PDF

FIT

Authors	Jaewon Chu et al.
Affiliations	Korea University; KAIST; Meta AI; UNIST
Central question	How can a multi-agent system identify which agent caused a failure and turn that evidence into useful prompt updates?
Method	Sequentially intervene on one agent, derive agent-level textual gradients, cluster similar gradients, and abstract a shared correction.
Key findings — paper claim	Reports state-of-the-art performance on five MAS benchmarks and 2.5× faster optimization than the next baseline.
Limitations	No public repository was verified today; prompt-only changes may not address tool, data, or runtime faults.
Why it matters — interpretation	Useful optimization layer for AutoNateAI orchestration, routing, and agent-team diagnostics.
Code / project	No verified repository

#6 SAEScientist-Bench: Can AI Agents Conduct Autonomous SAE Interpretability Research?

89

Submitted 8 Sep 2026 | cs.AI / cs.CL / cs.LG | PDF

FIT

Authors	Yuqiao Tan, Shizhu He, Jun Zhao, Kang Liu
Affiliations	CAS Institute of Automation; University of Chinese Academy of Sciences
Central question	Can AI agents autonomously perform sparse-autoencoder interpretability research rather than isolated analysis tasks?
Method	Ten agent configurations tackle 20 research tasks using Gemma Scope features for Gemma-2-9B-IT; evaluation covers discovery, causal steering, and measurement interpretation.
Key findings — paper claim	Frontier agents display discovery ability but remain behind experts, with causal steering and measurement interpretation as core weaknesses.
Limitations	Work-in-progress preprint; one model/SAE ecosystem and 20 tasks limit generalization.
Why it matters — interpretation	A strong model for benchmarking research agents and for adding measurement-literacy controls to AutoNateAI research workflows.
Code / project	Open repository

Paper briefs // 07-08

#7 Graph-Based Personalized Memory for LLM Agents

88

Submitted 8 Sep 2026 | [cs.AI](#) | [PDF](#)

FIT

Authors	Dac Duy Anh Nguyen, Zhangchi Qiu, Shigeng Chen, Alan Wee-Chung Liew
Affiliations	Griffith University
Central question	What design choices define graph-based personalized memory from representation through evaluation?
Method	Survey and taxonomy spanning representation, evolution, retrieval, and evaluation, with attention to explicit relations, time, provenance, and evidence links.
Key findings — paper claim	Synthesizes how graph memory models connected, revisable user knowledge and identifies evaluation gaps for long histories and changing preferences.
Limitations	Survey evidence depends on underlying studies; it does not establish a single best architecture. Accepted at ICKG 2026.
Why it matters — interpretation	Provides a design checklist for AutoNateAI's knowledge graph and personal-context layer.
Code / project	No paper-specific repository

#8 Time-Varying Data as Sheaves: an Invitation to Narratives

86

Submitted 8 Sep 2026 | [cs.AI](#) / [cs.MA](#) / [eess.SY](#) / [math.CT](#) | [PDF](#)

FIT

Authors	Wilmer Leal et al.
Affiliations	University of Florida; USP; Hamburg; National University of Colombia; Virginia Tech
Central question	How can changing data and switching relationships be represented without losing structural meaning?
Method	Book-chapter framework uses sheaves and 'narratives' to model time-varying objects, invariants, structural decomposition, and switching multi-agent topology.
Key findings — paper claim	Offers a unifying formal language rather than a benchmark result; examples connect dynamic data to networked systems and multi-agent communication.
Limitations	Highly mathematical and exploratory; operational tooling, empirical comparisons, and developer-facing implementation are future work.
Why it matters — interpretation	Potential theoretical backbone for Nathan's process philosophy, temporal graphs, simulations, and GIS change models.
Code / project	No verified repository

Paper briefs // 09-10

#9 Style Over Substance: Content-Invariant Wrappers Flip LLM Safety-Judge Verdicts

85

Submitted 8 Sep 2026 | cs.AI | PDF

FIT

Authors	Yongxi Zhou, Wenbo Ye, Yuanzhe Liu, Zihan Dong, Junwei Yao
Affiliations	Northeastern; Georgia Tech; USC
Central question	Can stylistic wrappers change automated safety verdicts even when the substantive content is fixed?
Method	Apply up to seven wrappers to 600 JailbreakBench replies and score them with eight judges; validate content invariance and judge errors with humans.
Key findings — paper claim	Reports a 19.9% unsafe-verdict flip rate for a GPT-4o-mini judge under a token refusal wrapper versus a 0.5% noise floor; some guards are far more robust.
Limitations	Limited wrapper family, judges, and safety dataset; agent-trajectory judging remains future work.
Why it matters — interpretation	Direct stress test for automated policy evaluation, governance dashboards, and ensemble judge design.
Code / project	Open repository

#10 Revoked but Still Authoritative: Revocation Enforcement in Agent-Memory Systems

84

Submitted 8 Sep 2026 | cs.AI / cs.CR | PDF

FIT

Authors	Yi Ting Shen, Kentaroh Toyoda, Alex Leung
Affiliations	Vulcan Research, AIFT
Central question	Do persistent-memory systems actually prevent revoked facts from being retrieved and acted upon?
Method	Load five systems with an obsolete policy and replacement; test nine policy scenarios, nine models, and six defenses while tracking retrieval and downstream action.
Key findings — paper claim	Paper reports that no tested system enforces revocation by default and that visible revocation labels do not reliably prevent obsolete facts from dominating retrieval/action.
Limitations	Five systems and constructed policy scenarios cannot cover all production memory architectures; defense durability under adaptive attacks is uncertain.
Why it matters — interpretation	Urgent adversarial complement to McClear and graph-memory design; ideal for a revocation benchmark.
Code / project	Open repository

Paper-to-code crosswalk

Paper	Repository	Reproduction target	Extension target
ExecCritic	MSR-Orchard/execcritic	Small SWE task; same-agent vs independent tester	Fail-closed customer CI gate
PlannerForge	TUM-AVS/PlannerForge	Generate/retrieve/mutate scenarios	Logistics, GIS, facilities simulation
MeClear	FudanSELab/MeClear	Recall and recovery on memory conflicts	Revocation-aware graph retriever
SAEScientist	Trae1ounG/SAEScientist	Run one interpretability task	Measurement-critic agent
Style Over Substance	safety-judge-robustness	Wrapper flip-rate baseline	Ensemble/rubric robustness dashboard
Revoked but Still Authoritative	Memory-Rebirth-Attack	Reproduce one system/model matrix	Adaptive revocation defense benchmark


Repository triage rule

Before adopting results, pin versions, inspect licenses and data dependencies, run a minimal happy-path example, then reproduce one headline metric. Only after the metric is understood should the code be adapted into an AutoNateAI system.

Code-availability gap

No paper-linked public repository was verified for Procedural Graphs, AgentGrad, or the Narratives chapter. Their proposed experiments therefore begin with a clean-room prototype or conceptual translation, not a claim of reproduction.

Authors and labs to know



Person / lab	Public profile	Why connect
Yuxing Lu	arXiv author search	Procedural graph design and evaluation protocol
Leitian Tao	Personal profile	Agentic post-training and coding-agent self-reflection
Jianfeng Gao	Microsoft Research	ExecCritic senior collaborator; general AI systems
Yuan Gao	TUM profile	PlannerForge implementation and scenario pipeline
Johannes Betz / TUM AVS	Lab	Agent-driven simulation and autonomous-systems validation
Shizhu He / NLPR	Paper contact	Autonomous interpretability-research evaluation
Vulcan Research, AIFT	Paper/contact	Memory revocation attacks and defense design

Contact discipline

Lead with a specific replication or extension question. Do not send generic admiration. The most credible opening is a one-paragraph experiment plan plus a clear statement of what result or artifact AutoNateAI can contribute.

Experiment queue // candidates only

These are not Airtable Experiments or Projects. They remain candidate ideas until Nathan selects or starts one.

Priority	Candidate experiment	Source papers	Core measurement	Expected artifact
1	Procedural Graph Runtime	Procedural Graphs + AgentGrad	Completion, wrong-order tools, repeats, latency, token cost	Graph runtime + benchmark report
2	Independent Test Agent	ExecCritic	Resolved tasks; false confidence; test validity; compute	Fail-closed repo-repair harness
3	Revocable Memory Graph	MeClear + Revocation + Graph Memory	Revoked retrieval/action, recovery, collateral forgetting	Memory governance test suite
4	Agentic Simulation Foundry	PlannerForge	Scenario validity, diversity, retrieval, edit success, system failures	Domain-neutral simulation pipeline
5	Safety Judge Stress Lab	Style Over Substance	Flip rate, noise floor, ensemble agreement, rubric sensitivity	Judge robustness dashboard
6	Research-Agent Measurement Critic	SAEScientist-Bench	Interpretation error, causal claims, reproducibility	Benchmark extension + critic agent
7	Temporal GIS Narratives	Narratives + graph-memory survey	Change consistency, provenance, query fidelity	Temporal graph / map prototype

Fastest start

Start Experiment 2 if Nathan wants a near-term developer-tool artifact with public code. Start Experiment 3 if the goal is a distinctive research thesis. Start Experiment 4 if the goal is a visually compelling simulation product that can cross into logistics, GIS, robotics, or infrastructure consulting.

Publication ideas // 01-04

1. Self-Evolving Procedural Graphs for Production Agents

Sources: Procedural Graphs + AgentGrad

Thesis/question: Can validation-gated graph edits improve long-horizon business automation without increasing unsafe tool use?

Section outline: 1) Problem and failure taxonomy; 2) graph representation; 3) localization/guidance; 4) intervention-based updates; 5) evaluation; 6) deployment lessons

Method/simulation: Build static and self-evolving graph controllers; compare with flat-history and text-memory baselines on 30-50 tool workflows.

Code to investigate: Clean-room graph runtime; investigate no official code yet.

Expected artifact: Repository, benchmark dataset, interactive trace viewer, technical paper.

Outreach: Ask Yuxing Lu about graph-edit validation and AgentGrad authors about intervention granularity.

2. Independent Behavioral Tests for Coding Agents

Sources: ExecCritic

Thesis/question: When does a separate test agent improve repository repair, and when do generated tests create false confidence?

Section outline: 1) Behavioral contracts; 2) role separation; 3) harness; 4) test validity; 5) repair results; 6) cost/security

Method/simulation: Run same-model same-trajectory, independent-agent, and human-test conditions on small open repositories; audit false positives and leaked information.

Code to investigate: MSR-Orchard/execcritic; SWE-agent mini harness.

Expected artifact: Reproduction notebook, secure harness, CI integration guide, blog/paper.

Outreach: Send Leitian Tao a compact replication matrix and ask which test-quality failure mode is most underexplored.

3. Revocable Personal Memory Graphs

Sources: MeClear + Revoked but Still Authoritative + Graph Memory survey

Thesis/question: Can a graph retriever enforce revocation structurally while preserving useful personalized memory?

Section outline: 1) Threat model; 2) temporal/provenance graph; 3) clearance attribution; 4) attack suite; 5) defenses; 6) UX/governance

Method/simulation: Combine soft-deletion, validity intervals, authorization filters, and MeClear attribution; test obsolete-policy retrieval and action across models.

Code to investigate: FudanSELab/MeClear + VulcanLab/Memory-Rebirth-Attack.

Expected artifact: Open benchmark, reference retriever, red-team report, architecture note.

Outreach: Contact both defense and attack teams with a proposed shared interoperability benchmark.

4. Agentic Scenario Testing Beyond Autonomous Driving

Sources: PlannerForge

Thesis/question: Does the generate→retrieve→mutate→execute→score pipeline transfer to logistics, GIS, facility operations, or infrastructure planning?

Section outline: 1) Transfer hypothesis; 2) scenario schema; 3) tool adapters; 4) retrieval/mutation; 5) simulation; 6) generalization

Method/simulation: Choose one domain, create 100 scenarios, compare agent and template generation, measure validity/diversity and discovered system failures.

Code to investigate: TUM-AVS/PlannerForge; relevant domain simulator.

Expected artifact: Scenario corpus, adapter SDK, evaluation dashboard, applied-research article.

Outreach: Ask Yuan Gao/Johannes Betz which PlannerForge modules are intentionally domain-independent.

Publication ideas // 05-07

5. How Robust Is Your AI Judge?

Sources: Style Over Substance

Thesis/question: Which content-invariant transformations produce the highest disagreement across safety judges, and can simple ensembles resist them?

Section outline: 1) Judge architecture; 2) wrapper taxonomy; 3) experiment design; 4) flip analysis; 5) mitigations; 6) product implications

Method/simulation: Reproduce selected wrappers across open judges and policy prompts; compare single judge, majority vote, alternate rubric, and deterministic checks.

Code to investigate: Yongxi-Zhou/safety-judge-robustness.

Expected artifact: Reproducible dashboard, wrapper corpus, mitigation checklist, publication.

Outreach: Share cross-domain wrapper results with Yongxi Zhou; ask about extending to agent trajectories.

6. Measurement Literacy for Autonomous Research Agents

Sources: SAEScientist-Bench

Thesis/question: Can a dedicated measurement critic reduce causal overclaiming and interpretation errors in autonomous AI research?

Section outline: 1) Research-agent workflow; 2) error taxonomy; 3) critic design; 4) benchmark extension; 5) ablation; 6) research governance

Method/simulation: Insert a critic that checks metric definition, uncertainty, causal intervention, and replication before a research agent finalizes a finding.

Code to investigate: Trae1ounG/SAEScientist.

Expected artifact: Benchmark extension, critic prompts/tools, scored trajectories, workshop paper.

Outreach: Ask Shizhu He which measurement-interpretation errors recur across agent configurations.

7. Temporal Narratives for Evolving GIS Systems

Sources: Time-Varying Data as Sheaves + Graph Memory survey

Thesis/question: Can narrative/sheaf structure preserve provenance and consistency in evolving geographic knowledge graphs better than event logs alone?

Section outline: 1) GIS change problem; 2) narrative formalism; 3) graph schema; 4) temporal queries; 5) simulation; 6) visualization

Method/simulation: Simulate parcel, infrastructure, or logistics-network changes; compare event sourcing, temporal property graphs, and a narrative-inspired representation.

Code to investigate: No paper-specific code; use NetworkX/Neo4j plus a mapping stack.

Expected artifact: Interactive temporal map, query benchmark, explainer, technical note.

Outreach: Contact Wilmer Leal with a concrete GIS translation diagram and questions about implementable invariants.

Portfolio strategy

The seven ideas form a coherent publication program: execution structure, independent verification, governed memory, simulation testing, robust evaluation, research-agent methodology, and temporal graph design. The common thesis is that capable agents need inspectable structure around generation.

Actions today

Priority	Owner / target	Action	Deliverable by end of day
1	Nathan	Deep-read Procedural Graphs; redraw the algorithm against one AutoNateAI workflow.	One-page graph schema + failure taxonomy
2	Nathan / ExecCritic repo	Run setup/README audit and choose one tractable repository task.	Reproduction plan with compute/security constraints
3	Nathan / TUM AVS	Inspect PlannerForge scenario representation and adapters.	Domain-transfer map for logistics or GIS
4	Nathan	Pair-read McClear and Revocation Enforcement.	Revocable-memory benchmark matrix
5	Yuxing Lu or Leitian Tao	Send one specific technical question only after completing the relevant map/replication plan.	Evidence-based outreach message

Today's decision gate



Airtable update

Created 10 Research Paper source records, 6 GitHub Repository source records, 5 People records, and 3 Companies/labs records after canonical-ID/URL/name deduplication. No Lab Publications, Experiments, or Projects were created. The current Sources schema has no linked-record field to People or Companies, so relationships were preserved in author/lab notes rather than fabricating links.

Methodology, evidence, and source index

Method

The scan covered arXiv recent submissions and adjacent computer-science categories, then prioritized papers by AutoNateAI relevance, system novelty, evidence clarity, code availability, and experimentability. Each selected title, author list, date, category, affiliation, and code claim was checked against the arXiv abstract page and/or paper PDF. Paper claims and AutoNateAI interpretations are explicitly labeled in briefs.

Limitations of this edition

A one-day radar cannot be literally exhaustive across every repository, venue, and discussion channel. arXiv abstracts and author PDFs remain preliminary unless a venue acceptance is explicitly reported. Absence of a code link means 'not verified today,' not proof that no code exists. Citations and social discussion are too immature for today's submissions to serve as quality signals.

Paper / source	Abstract	PDF / code
Procedural Graphs: Self-Evolving Execution Structures for LLM Agents	2609.09153	PDF
ExecCritic: Learn to Test, Test to Improve for Coding Agents	2609.09133	PDF · Code
PlannerForge: LLM Agents for Scenario-Based Testing of Motion Planners	2609.08965	PDF · Code
McClear: Cooperative Game-Theoretic Attribution and Risk-Aware Memory Clearance	2609.09115	PDF · Code
AgentGrad: Intervention-guided Prompt Optimization for Multi Agent Systems	2609.08572	PDF
SAEScientist-Bench: Can AI Agents Conduct Autonomous SAE Interpretability Research?	2609.09113	PDF · Code
Graph-Based Personalized Memory for LLM Agents	2609.08599	PDF
Time-Varying Data as Sheaves: an Invitation to Narratives	2609.09056	PDF
Style Over Substance: Content-Invariant Wrappers Flip LLM Safety-Judge Verdicts	2609.08236	PDF · Code
Revoked but Still Authoritative: Revocation Enforcement in Agent-Memory Systems	2609.08258	PDF · Code

Prepared for Nathan Baker / AutoNateAI. Research intelligence is not peer-review, replication, or endorsement. Preserve source/version identifiers when reusing claims.

Reusable Mermaid source

```
graph LR
    A[Paper] --> B[Code]
    B --> C[Experiment]
    C --> D[Artifact]
    D --> E[Publication]
```